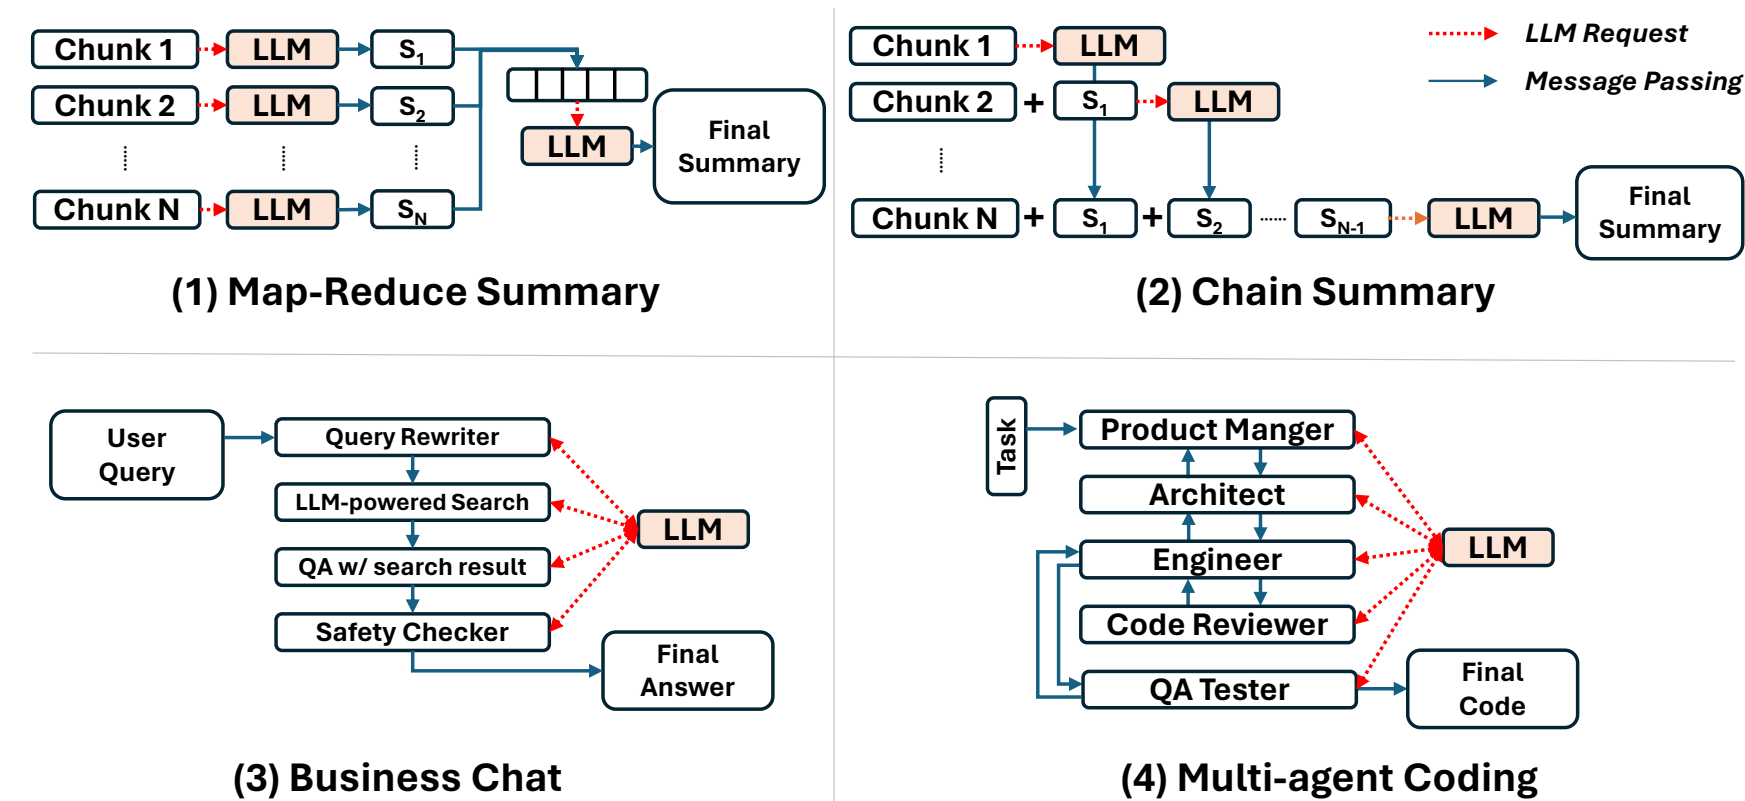
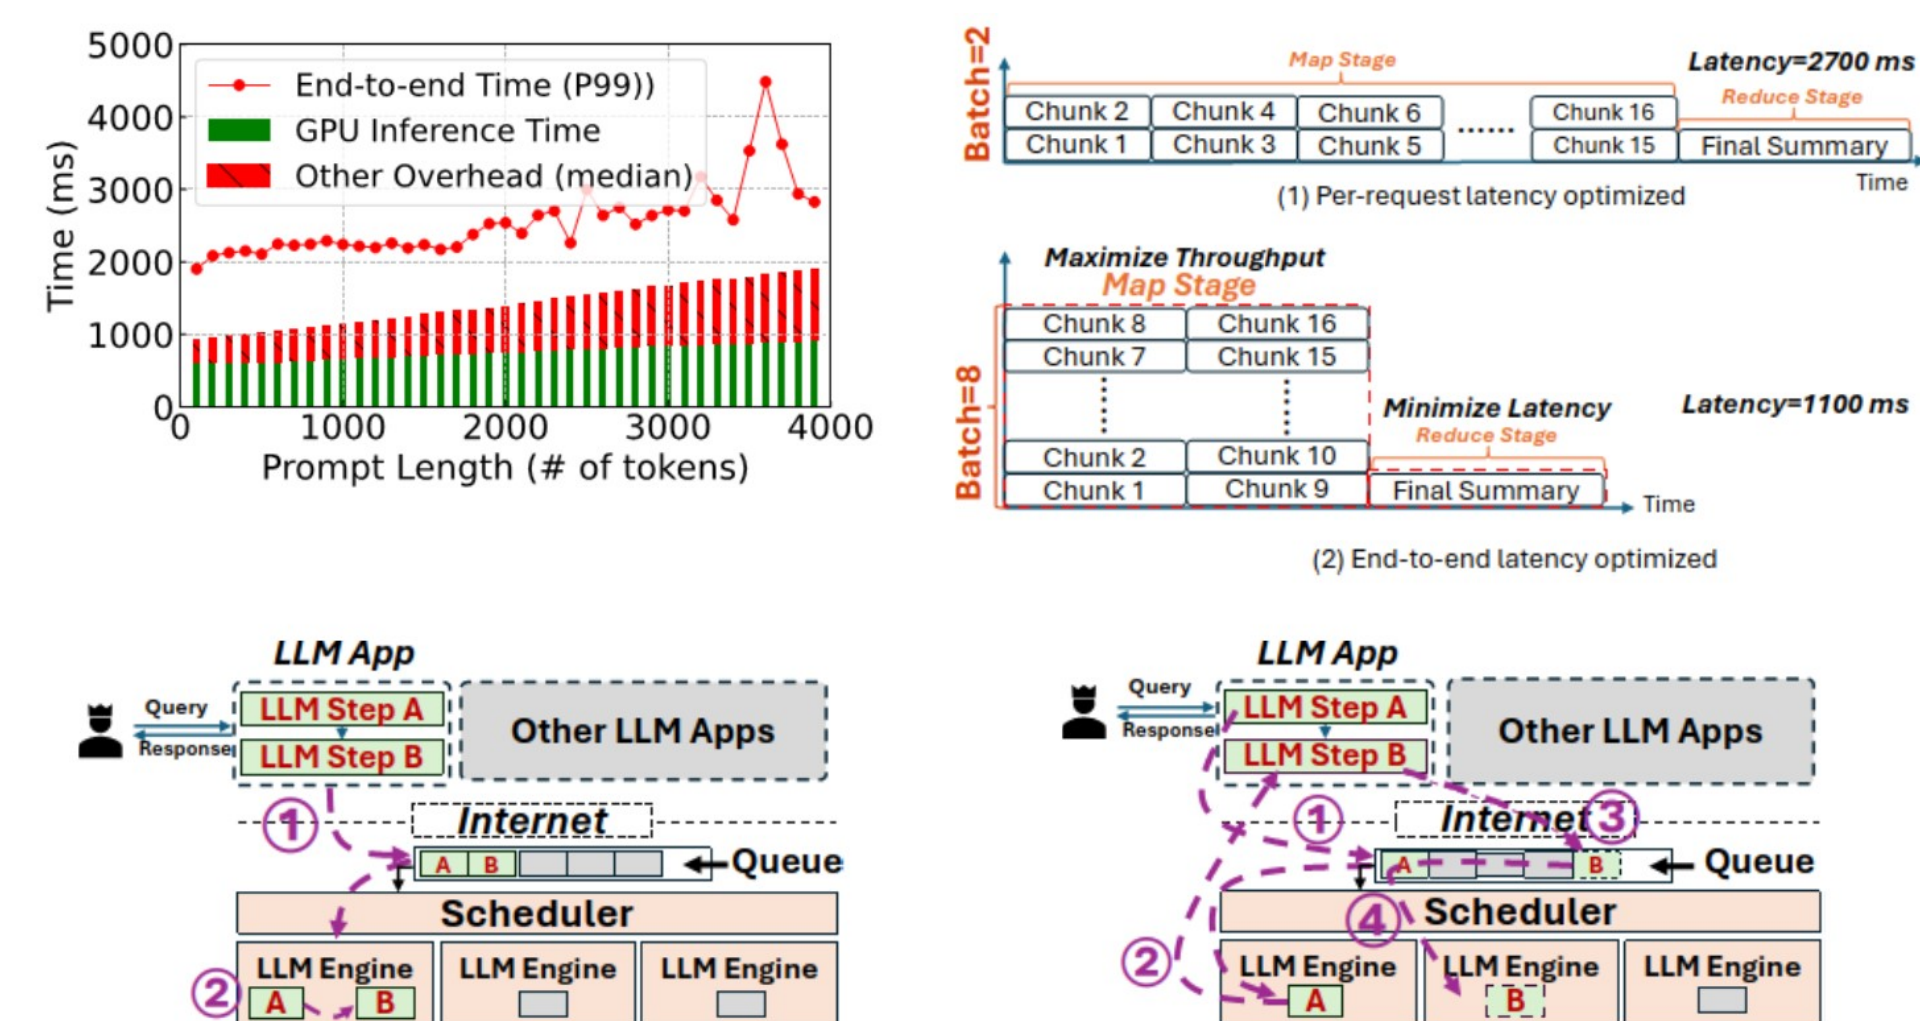


Multi-tenant LLM Services face Diverse Applications (LLM Apps)



Public multi-tenant LLM services use too over-simplified request-level API, which **loses the essential application-level information**, leading to sub-optimal end-to-end performance.

Problems of Multi-tenant Serving



- (1) Excessive Overhead of Consecutive Requests;
- (2) Misaligned Scheduling Objectives;
- (3) Redundant Computations

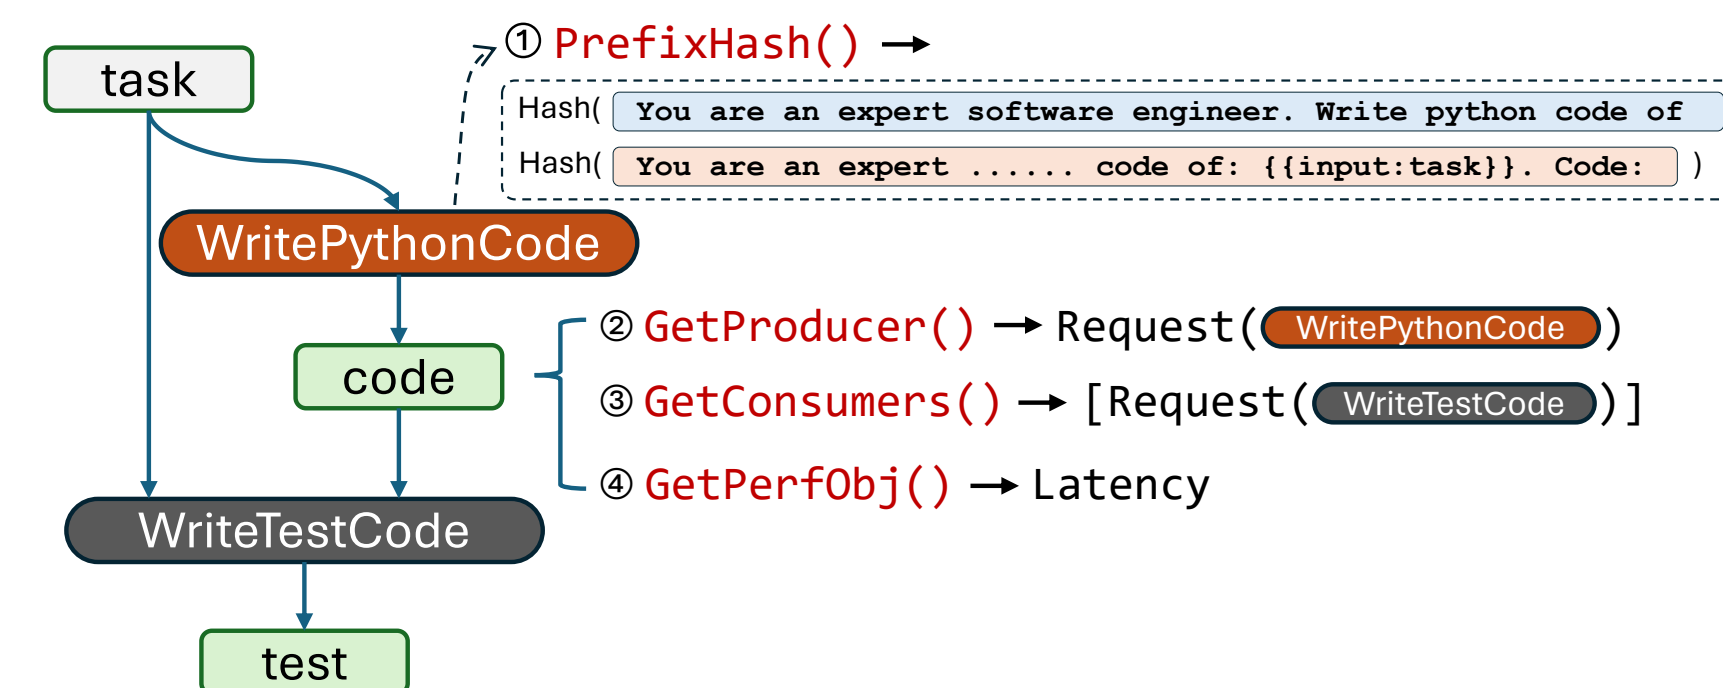
Parrot Design with Semantic Variable as Core Abstraction

```
import Parrot as P
from Parrot.PerformanceCriteria import LATENCY

@P.SemanticFunction
def WritePythonCode(task: P.SemanticVariable):
    """ You are an expert software engineer.
    Write python code of {{input:task}}.
    Code: {{output:code}}
    """

@P.SemanticFunction
def WriteTestCode(
    task: P.SemanticVariable,
    code: P.SemanticVariable):
    """ You are an experienced QA engineer.
    You write test code for {{input:task}}.
    Code: {{input:code}}.
    Your test code: {{output:test}}
    """

def WriteSnakeGame():
    task = P.SemanticVariable("a snake game")
    code = WritePythonCode(task)
    test = WriteTestCode(task, code)
    return code.get(perf=LATENCY), test.get(perf=LATENCY)
```

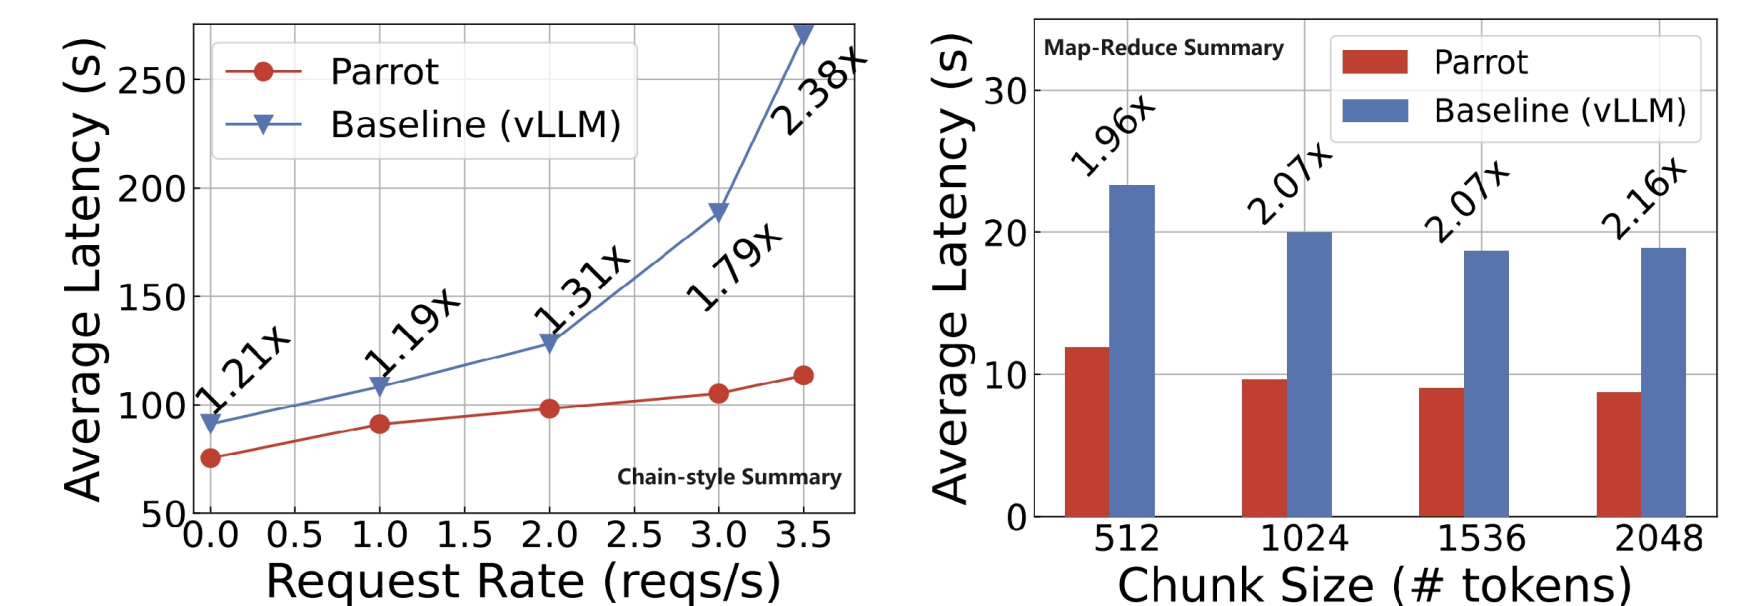


Parrot provides a natural way of **programming LLM applications with Semantic Variable annotations**, enabling DAG analysis and optimizations.

Optimizations with Semantic Variable


1. Serving Dependent Requests.

Serving dependent requests. Avoid unnecessary communication. Optimized scheduling.



2. Sharing Prompt Prefix.

Quickly detect (static & dynamic) sharing, maximize them by scheduling, run them with efficient kernels.



3. App-Centric Scheduling.

Mixed workloads of latency and throughput preference. Parrot achieves the best of both worlds.

